


The Company & its Products

| Bloomberg Terminal Demo Request

| [Bloomberg Anywhere Login](#)

| Customer Support

[Subscribe](#)



• **LIVE**


[How to Navigate the 'AI Scare Trade': Live Q&A](#)

[Join Now](#)

[Newsletter](#) | Odd Lots

Presenting FedLock, A New Way to Measure FedSpeak With LLMs

Check it out.



By [Joe Weisenthal](#) and [Tracy Alloway](#)

March 5, 2026 at 4:39 PM UTC



Save



Translate



This article is for **subscribers only**.



You're reading the Odd Lots newsletter.

Subscribe now to get Joe and Tracy's daily newsletter on the newest market crazes.

[Subscribe](#)

Hello and welcome to the newsletter, a grab bag of daily content from the Odd Lots universe. Sometimes it's us, Joe Weisenthal and Tracy Alloway, bringing you our thoughts on the most recent developments in markets, finance and the economy. And sometimes it's contributions from our network of expert guests and sources. Whatever it is, we promise it will always be interesting.

If you like chatting with us, check out the [Odd Lots Discord](#), where you can hang out and talk with us and with other listeners 24/7.

What Joe is thinking about today

Measuring sentiment across large bodies of text is obviously something of high interest to many businesses and investors, and nowhere is this more true than in monetary policy. There's a whole field of Fed Kremlinology, in which economists scrutinize the utterances of central bankers to glean clues around the future direction of interest rates. And so, as an experiment to learn more about how LLMs can be used for this type of thing, I have a new hobby project called [FedLock](#). Check it out. And then after you check it out read the below.

Now, some background. Earlier in the year [I wrote](#) about my first vibecoding project called [Havelock](#), which is an attempt to build software that can detect whether text more closely resembles the written or spoken word, using a lexical framework described in Walter Ong's book *Orality and Literacy*.

I'm not going to recount everything I've done so far for Havelock, but the basic gist is that I needed a rudimentary model that can recognize where on the orality-literacy spectrum a given text exists. The only real hard part in that is giving the model data to train on.

Let's say you were creating a model that could recognize the difference between a dog and a cat. Well, intuitively you'd start by uploading thousands of cat pictures and thousands of dog pictures, so that the model can learn the contours of each. My problem is that (as far as I know) there's no existing dataset that divides texts along the Walter Ong framework.

So I had to create my own training data. My first attempt at doing this was to upload thousand of texts from sites like Project Gutenberg and Wikipedia Talk pages and others, and ask an LLM (Claude Sonnet in this case) to grade them on a range of one to 100, with 100 being the most oral and one being the most literate.

The problem I encountered is that almost everything ended up being either 80 or 20 (more on this later), and I was having trouble really getting the full spectrum. Hardly anything was falling in the middle. Eventually I landed on a different approach. Rather than asking for a direct score, I built a tournament, where the LLM would be presented with two different randomly selected texts in my corpus, and then was then simply asked basically something along the lines of: "Of these two texts, which of them is more oral, as understood by the Walter Ong framework." After presenting the LLM thousands of those pairwise comparisons, I was able to generate a nice spectrum in which things like scientific paper abstracts and legal briefs all score highly literate, random Substacks


score more in the middle, and everything from Homer to standup comedy scores as highly oral. With these graded texts, I was then able to go on and train the actual model itself to recognize the difference. (For what it's worth, there's still tons of work that needs to be done on Havelock, especially around identifying specific lexical markers but basically it works as intended. You can try it out.)

Anyway, after building that, I started wondering whether there's a way to repurpose the model to central bank speak. Perhaps I could detect interesting things about whether central bankers communicate in a more oral register during times of crisis, or something like that. So far I haven't actually determined very much on that front, [but I have had some interesting findings](#), such as the fact that Fed district presidents tend to be more oral in their communications than, say, Fed chairs or governors.

But then it occurred to me that the basic tournament model I used to generate training data might be interesting as a simple way to measure where Fed communication stands on the well known hawk-dove spectrum.

And so I went about, basically, trying to find every speech given by an FOMC member over the last 30 years, removing the ones that weren't about monetary policy and then just asking the LLM (in this case Gemini Flash) to determine which of the two were more hawkish, given the conditions of the time.

It looks like this:



Source: FedLock

The LLM judge sees two speeches side by side, with the name of the speaker removed (so that it's not biased by what it knows about the speaker's hawkish/dovish biases), and with the economic data at the time of the speech.

After doing these pairings 65,000 times across the entire corpus, a

whole spectrum emerges from the least to most hawkish speech. Then that spectrum is converted to a one to 100 scale and voilà.

Once every speech is scored, it then becomes possible to average them together (more or less) into a time series that captures FOMC-wide sentiment as a whole:



It looks pretty good, I'd say. There's no way to prove that this exact measure correctly gets the sentiment of the FOMC at any given time. But if you overlay it with something like the Taylor Rule, I'd say it tracks pretty well with general shifts.



Now, a few notes:

- If you go to [FedLock](#), you can read the whole methodology (which, to be clear, was written by Claude itself, with a little input from me. There's a lot more in there).
- 65,000 pairwise comparisons with Gemini Flash cost about \$33 and took about 15 minutes. It's kind of shocking how much you can do with relatively low-end LLMs. I did some other tournaments using other models, including the (much) more expensive Claude Sonnet and also DeepSeek v3, and the results looked almost the same, so it's clear that you don't need a state-of-the-art model for this kind of analysis.
- I mentioned above that one of the problems with just asking the AI to score something on a one to 100 scale, is that the distribution tends to be bimodal, and this was the case with asking about a hawkish-dovish scale. Here you can see the difference in scores/results between the two methods:

- One of the main potential issues that arises in this type of thing is that all of the speeches given by FOMC members are theoretically in the training data itself. And the AI may be biased to assume that certain speakers are doves or hawks, and this may influence the judgments. I ran a number of tests, including putting fake or different names on the speakers. In one test, I replaced the well-known hawk Jeffrey Lacker with an imaginary Regional Fed President Whitfield, and the speeches still scored roughly the same as in the anonymized test. You can see in the methodology that name identification did have a measurable effect, but that generally speaking, there's good reason to think that the LLM is actually judging the text as given at the time, in the conditions at the time.
- This test only uses Fed speeches and press conferences. None of the other modes of communication (such as TV interviews) are in this, so some signal is being missed. That's just a matter of assembling a more complete corpus.

Anyway, the main point of this project was not to learn more about the Fed, it was to learn more about the capabilities of LLMs.

With this approach, there are probably numerous other things you could test out in this way. For example, if you had earnings call transcripts, you could rank them all on management confidence, or what they're saying about pricing power, or what they're saying about aggregate demand, and then potentially derive signal from that.

At some point I want to do all that. But in the meantime, [I'd love it if you checked out FedLock](#) and tell me what does and doesn't resonate, and how it could be improved. Because it's so cheap and easy and I have a (mostly) complete corpus, I can run plenty of other tests. And each time a new speech comes out, it can be inserted into the tournament for a cost of about a penny or two.

On the podcast

Today we had the pleasure of speaking with **Lloyd Blankfein, the former CEO of Goldman Sachs**. Lloyd has been relatively quiet since retiring in 2018. He didn't go into government. He's not running some new private credit vehicle. But he is out with a new book called [Streetwise: Getting to and Through Goldman Sachs](#), about his life and career. And so this was the perfect hook to talk with him at the recent Bloomberg Invest conference. We discussed everything from private credit to bailouts to the future of NYC, as well as how he resists the impulse to tweet.

Listen: [Apple](#) / [Spotify](#)

More from Bloomberg

Enjoying Odd Lots? Check out these newsletters:

- [Markets Daily](#) for what's moving in stocks, bonds, FX and commodities
- [Going Private](#) for coverage of private markets and the forces moving capital away from the public eye
- [CFO Briefing](#) for what finance leaders need to know
- [Deals](#) for transactions and analysis, from IPOs to startup investing
- [Money Stuff](#) for Bloomberg Opinion's Matt Levine's daily newsletter on Wall Street and finance

You have exclusive access to other subscriber-only newsletters. [Explore all newsletters](#) here to get most out of your Bloomberg subscription.

Get Alerts for:

+ Joe Weisenthal

+ Tracy Alloway



Save

Contact us:

[Provide news feedback or report an error](#)

Site feedback:

[Take our Survey](#)

Confidential tip?

[Send a tip to our reporters](#)

Before it's here, it's on the [Bloomberg Terminal](#)

More From Bloomberg



US Sinking of Iranian Warship Piles Pressure on India's Modi

updated an hour ago




China Tells Top Refiners to Halt Diesel and Gasoline Exports



Apple Rolls Out \$599 MacBook Neo, Threatening Windows PC Market



UAE and Qatar Urge Allies to Help Trump Find Iran Off-Ramp



Sony Pulls Back From PlayStation Games on PC

Top Reads



Iran Satellite Images Reveal Trump's Strategy Risks Chaotic Endgame

by Peter Martin, Krishna Karra, Jonathan Tirone, Raedeh Wahid, Dave Merrill and Tom Fevrier



Iran War Oil Shock Threatens to Unleash Wave of Global Inflation

by Ziad Daoud, Dina Esfandiary, Jamie Rush, Jennifer Welch and Tom Orlik, Bloomberg Economics



The Caribbean Nation Unexpectedly Going Along With Trump's Drug-Boat Narrative

by Michael Moore



Trump's War on Iran Has Traders Staring Down an Energy Crisis

by Jack Farchy, Alex Longley, Priscila Azevedo Rocha, Charles Gorrivan and Julian Lee

[Home](#)

[BTV+](#)

[Market Data](#)

[Opinion](#)

[Audio](#)

[Originals](#)

[Magazine](#)

[Events](#)

News

[Markets](#)

[Economics](#)

[Technology](#)

[Politics](#)

[Green](#)

[Crypto](#)

[AI](#)

[Terms of Service](#)

[Manage Cookies](#)

[Trademarks](#)

[Privacy Policy](#)

[Careers](#) [Advertise](#) [Ad Choices](#) [Help](#)

©2026 Bloomberg L.P. All Rights Reserved.

